

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
```

```
In [2]: d = r'C:\Users\Pratik Sonawane\Downloads\diabetes.csv'
df = pd.read_csv(d)
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [4]: df.shape
```

```
Out[4]: (768, 9)
```

```
In [5]: df.isnull().sum()
```

```
Out[5]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness 0
Insulin      0
BMI          0
DiabetesPedigreeFunction 0
Age          0
Outcome      0
dtype: int64
```

```
In [6]: df.describe()
```

```
Out[6]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

In [7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [8]: *# Target Variababl*

```
df['Outcome'].value_counts()
```

Out[8]: Outcome
0 500
1 268
Name: count, dtype: int64

- 500 patients (about 65%) have Outcome = 0 (no diabetes).
- 268 patients (about 35%) have Outcome = 1 (diabetes).

```
In [9]: grouped_means = df.groupby('Outcome').mean().T

# Calculate percentage difference
grouped_means['% diff'] = 100 * (grouped_means[1] - grouped_means[0]) / grouped_means[0]

print(grouped_means)
```

Outcome	0	1	% diff
Pregnancies	3.298000	4.865672	47.534010
Glucose	109.980000	141.257463	28.439228
BloodPressure	68.184000	70.824627	3.872795
SkinThickness	19.664000	22.164179	12.714499
Insulin	68.792000	100.335821	45.853909
BMI	30.304200	35.142537	15.965897
DiabetesPedigreeFunction	0.429734	0.550500	28.102501
Age	31.190000	37.067164	18.843104

High differences (above 25%):

- Pregnancies: Diabetic patients have almost 48% more pregnancies on average, suggesting a potential link between pre-pregnancy factors and diabetes risk.
- Glucose: Blood glucose levels are around 28% higher in the diabetic group, highlighting impaired glucose regulation as a major characteristic of the disease.
- Insulin: Diabetic patients show a significant 46% increase in average insulin levels, indicating insufficient or ineffective insulin production or utilization.
- DiabetesPedigreeFunction: A 28% increase in this genetic score suggests a possible hereditary influence on diabetes risk.

Moderate differences (around 10-25%):

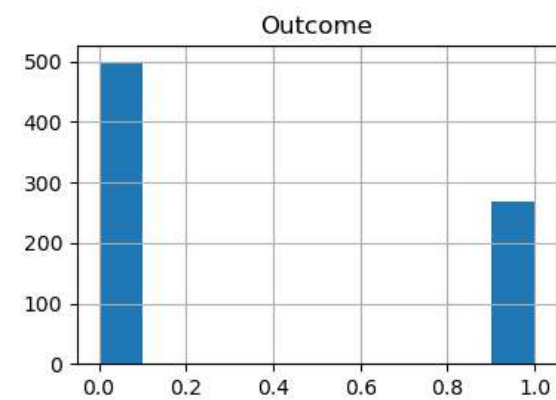




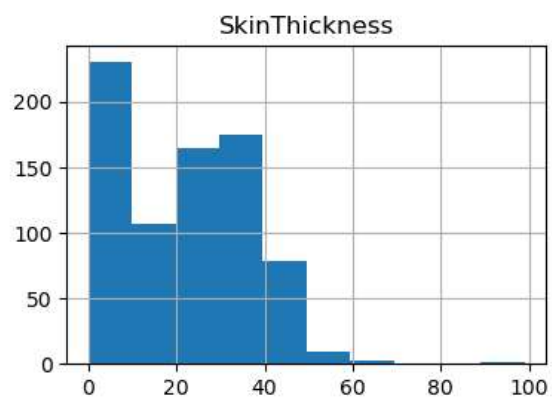


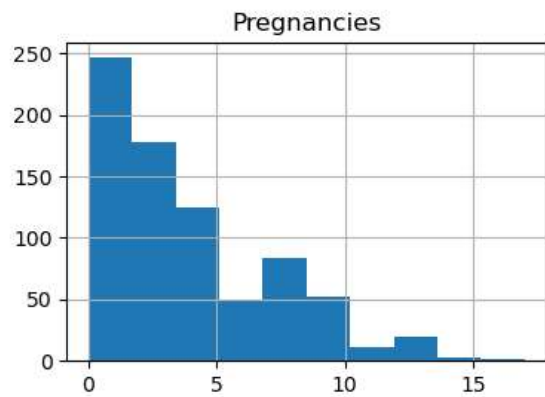
- BMI: Diabetic patients have a 16% higher average BMI, reflecting a potential role of body weight in diabetes development.
- SkinThickness: Increased skin thickness by 13% could be a marker for insulin resistance or metabolic changes associated with diabetes.
- Age: There's a 19% difference in average age, with diabetic patients being older on average. This could be due to age-related changes in metabolism or longer exposure to risk factors.

Low difference (below 5%):

- BloodPressure: While slightly higher in diabetic patients, the difference in average blood pressure is relatively small, indicating it might not be a primary factor for everyone.


```
In [10]: df.hist(bins=10,figsize = (15,10))
```

```
Out[10]: array([[<Axes: title={'center': 'Pregnancies'}>,
<Axes: title={'center': 'Glucose'}>,
<Axes: title={'center': 'BloodPressure'}>],
[<Axes: title={'center': 'SkinThickness'}>,
<Axes: title={'center': 'Insulin'}>,
<Axes: title={'center': 'BMI'}>],
[<Axes: title={'center': 'DiabetesPedigreeFunction'}>,
<Axes: title={'center': 'Age'}>,
<Axes: title={'center': 'Outcome'}>]], dtype=object)
```



```
In [11]: df.corr()['Outcome']
```

```
Out[11]: Pregnancies      0.221898  
Glucose      0.466581  
BloodPressure  0.065068  
SkinThickness 0.074752  
Insulin      0.130548  
BMI          0.292695  
DiabetesPedigreeFunction 0.173844  
Age          0.238356  
Outcome      1.000000  
Name: Outcome, dtype: float64
```

- A high correlation simply indicates a linear relationship between two variables, but it doesn't necessarily mean one directly causes the other.
- Features with seemingly weak correlations might still be valuable based on ones understanding of the problem and the underlying biology.

```
In [15]: X = df.drop(columns = ['Outcome'])  
Y = df['Outcome']
```

```
In [17]: scaler = StandardScaler()
```

```
In [19]: scaler.fit(X)
```

```
Out[19]: StandardScaler()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [20]: standard_data = scaler.transform(X)
```



```
In [22]: X_scaled = scaler.fit_transform(X)
X_scaled
```

```
Out[22]: array([[ 0.63994726,  0.84832379,  0.14964075, ...,  0.20401277,
                  0.46849198,  1.4259954 ],
                [-0.84488505, -1.12339636, -0.16054575, ..., -0.68442195,
                  -0.36506078, -0.19067191],
                [ 1.23388019,  1.94372388, -0.26394125, ..., -1.10325546,
                  0.60439732, -0.10558415],
                ...,
                [ 0.3429808 ,  0.00330087,  0.14964075, ..., -0.73518964,
                  -0.68519336, -0.27575966],
                [-0.84488505,  0.1597866 , -0.47073225, ..., -0.24020459,
                  -0.37110101,  1.17073215],
                [-0.84488505, -0.8730192 ,  0.04624525, ..., -0.20212881,
                  -0.47378505, -0.87137393]])
```

```
In [25]: X_train,X_test,y_train,y_test = train_test_split(X_scaled,Y,test_size = 0.2,stratify = Y,random_state=42)
```

```
In [27]: print(X.shape,X_train.shape,X_test.shape)
```

```
(768, 8) (614, 8) (154, 8)
```

```
In [28]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier
```

```
In [29]: lr = LogisticRegression()

        dtc = DecisionTreeClassifier()

        rf = RandomForestClassifier()

        svm = SVC()

        knn = KNeighborsClassifier()

        GBM = GradientBoostingClassifier()
```

```
In [31]: lr.fit(X_train,y_train)
        dtc.fit(X_train,y_train)
        rf.fit(X_train,y_train)
        svm.fit(X_train,y_train)
        knn.fit(X_train,y_train)
        GBM.fit(X_train,y_train)
```

Out[31]: GradientBoostingClassifier()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [32]: from sklearn.metrics import classification_report
```

```
In [34]: def model_evaluation(model):  
        y_pred = model.predict(X_test)  
        accuracy = accuracy_score(y_test, y_pred)  
        return y_pred, accuracy  
  
lr_pred, lr_accuracy = model_evaluation(lr)  
dtc_pred, dtc_accuracy = model_evaluation(dtc)  
rf_pred, rf_accuracy = model_evaluation(rf)  
svm_pred, svm_accuracy = model_evaluation(svm)  
knn_pred, knn_accuracy = model_evaluation(knn)  
GBM_pred, GBM_accuracy = model_evaluation(GBM)  
  
print(f"Logistic Regression accuracy: {lr_accuracy:.4f}")  
print(f"Decision Tree accuracy: {dtc_accuracy:.4f}")  
print(f"Random Forest accuracy: {rf_accuracy:.4f}")  
print(f"Support Vector Machine accuracy: {svm_accuracy:.4f}")  
print(f"K-Nearest Neighbors accuracy: {knn_accuracy:.4f}")  
print(f"Gradient Boosting Model accuracy: {GBM_accuracy:.4f}")
```

```
Logistic Regression accuracy: 0.7143  
Decision Tree accuracy: 0.7208  
Random Forest accuracy: 0.7532  
Support Vector Machine accuracy: 0.7468  
K-Nearest Neighbors accuracy: 0.7078  
Gradient Boosting Model accuracy: 0.7532
```

rf and GBM giving same score , preparing classification report and model with best recall score will be best for our project

```
In [39]: from sklearn.metrics import classification_report

rf_report = classification_report(y_test, rf.predict(X_test))
print('rf :',rf_report)

GBM_report = classification_report(y_test, GBM.predict(X_test))
print('GBM :',GBM_report)
```

```
rf :                precision    recall  f1-score   support

      0        0.79         0.84         0.82        100
      1        0.67         0.59         0.63         54

   accuracy                0.75         154
  macro avg        0.73         0.72         0.72         154
 weighted avg        0.75         0.75         0.75         154

GBM :                precision    recall  f1-score   support

      0        0.79         0.84         0.82        100
      1        0.67         0.59         0.63         54

   accuracy                0.75         154
  macro avg        0.73         0.72         0.72         154
 weighted avg        0.75         0.75         0.75         154
```

```
In [40]: from sklearn.model_selection import StratifiedKFold
         from sklearn.model_selection import cross_val_score
```

```
In [45]: lrcv = cross_val_score(lr,X_scaled,Y,cv = StratifiedKFold(n_splits=5)).mean()
print('lr :', lrcv)
dtccv = cross_val_score(dtc,X_scaled,Y,cv = StratifiedKFold(n_splits=5)).mean()
print('dtccv :', dtccv)
rfcv = cross_val_score(rf,X_scaled,Y,cv = StratifiedKFold(n_splits=5)).mean()
print('rfcv :', rfcv)
svmcv = cross_val_score(svm,X_scaled,Y,cv = StratifiedKFold(n_splits=5)).mean()
print('svmcv :', svmcv)
knncv = cross_val_score(knn,X_scaled,Y,cv = StratifiedKFold(n_splits=5)).mean()
print('knncv :', knncv)
GBMcV = cross_val_score(GBM,X_scaled,Y,cv = StratifiedKFold(n_splits=5)).mean()
print('GBMcV :', GBMcV)
```

```
lr : 0.7708853238265002
dtccv : 0.7201680672268906
rfcv : 0.7643833290892115
svmcv : 0.7708938120702827
knncv : 0.733112638994992
GBMcV : 0.7578728461081402
```

Predictive System using SVM Model

```
In [46]: X.columns
```

```
Out[46]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
               'BMI', 'DiabetesPedigreeFunction', 'Age'],
              dtype='object')
```

```
In [48]: Pregnancies = int(input('Enter the Pregnancies =',))
Glucose = float(input('Enter the Glucose =',))
BloodPressure = float(input('Enter the BloodPressure =',))
SkinThickness = float(input('Enter the SkinThickness =',))
Insulin = float(input('Enter the Insulin =',))
BMI = float(input('Enter the BMI =',))
DiabetesPedigreeFunction = float(input('Enter the DiabetesPedigreeFunction =',))
Age = float(input('Enter the Age =',))

patient_data = [Pregnancies,Glucose,BloodPressure,SkinThickness,Insulin,
                BMI,DiabetesPedigreeFunction,Age]
patient_data
```

```
Enter the Pregnancies =5
Enter the Glucose =120
Enter the BloodPressure =140
Enter the SkinThickness =28
Enter the Insulin =95
Enter the BMI =36
Enter the DiabetesPedigreeFunction =0.52
Enter the Age =42
```

```
Out[48]: [5, 120.0, 140.0, 28.0, 95.0, 36.0, 0.52, 42.0]
```

```
In [49]: input_array = np.asarray(patient_data)
input_array
```

```
Out[49]: array([ 5. , 120. , 140. , 28. , 95. , 36. , 0.52, 42. ])
```

```
In [50]: # Reshape in 2D
input_data = input_array.reshape(1,-1)
input_data
```

```
Out[50]: array([[ 5. , 120. , 140. , 28. , 95. , 36. , 0.52, 42. ]])
```

```
In [51]: std_data = scaler.transform(input_data)
std_data
```

C:\Users\Pratik Sonawane\anaconda3\Lib\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(

```
Out[51]: array([[ 0.3429808 , -0.02799627,  3.66508775,  0.4681735 ,  0.13198454,
                  0.50861896,  0.14533928,  0.74529338]])
```

```
In [52]: prediction = svm.predict(std_data)
prediction
```

```
Out[52]: array([0], dtype=int64)
```

```
In [53]: prediction[0]
```

```
Out[53]: 0
```

Model has accurately predicted above pateint data with no diabetes(0)

Prediction System

```
In [57]: import warnings
warnings.filterwarnings('ignore')

Pregnancies = int(input('Enter the Pregnancies =',))
Glucose = float(input('Enter the Glucose =',))
BloodPressure = float(input('Enter the BloodPressure =',))
SkinThickness = float(input('Enter the SkinThickness =',))
Insulin = float(input('Enter the Insulin =',))
BMI = float(input('Enter the BMI =',))
DiabetesPedigreeFunction = float(input('Enter the DiabetesPedigreeFunction =',))
Age = float(input('Enter the Age =',))

patient_data = [Pregnancies,Glucose,BloodPressure,SkinThickness,Insulin,
                BMI,DiabetesPedigreeFunction,Age]

input_array = np.asarray(patient_data)
input_data = input_array.reshape(1,-1)
std_data = scaler.transform(input_data)
prediction = svm.predict(std_data)

print('*****Model Predictions*****')
if (prediction[0] ==0):
    print('Result - Negative')
    print('The person is Not Diabetic')
else:
    print('Result - Positive')
    print('The Patient is Diabetic')
```



```
Enter the Pregnancies =6
Enter the Glucose =148
Enter the BloodPressure =72
Enter the SkinThickness =35
Enter the Insulin =0
Enter the BMI =33.6
Enter the DiabetesPedigreeFunction =0.627
Enter the Age =50
*****Model Predictions*****
Result - Positive
The Patient is Diabetic
```